

Reduce Before You Localize: Delta-Debugging and Spectrum-Based Fault Localization

Arpit Christi,* Matthew Lyle Olson,* Mohammad Amin Alipour,* Alex Groce*

Abstract

Debugging is one of the most time-consuming and difficult aspects of software development. Recent years have seen a wide variety of research efforts devoted to easing the burden of debugging by automatically localizing faults. The most popular approaches use spectra of failing and successful executions to score program entities according to how likely they are to be faulty. Since the original work on this topic, many formulas have been proposed to improve the accuracy of scores. Most of these improvements are either marginal or context-dependent. This paper proposes that, independent of the scoring method used, the effectiveness of spectrum-based localization can be dramatically improved by, when possible, delta-debugging failing test cases and basing localization only on the reduced test cases. Unlike most formula changes, under reasonable assumptions reduction can never reduce the effectiveness of localization; we show that in practice, reduction very seldom reduces localization effectiveness. Moreover, we show that for programs and faults taken from the standard localization literature, a large case study of Mozilla’s JavaScript engine using 10 real faults, and for mutants of various open source projects, localizing only after reduction often produces much better rankings for faults than localization without reduction, independent of the localization formula used, and the improvement is often even greater than that provided by changing from the worst to the best localization formula for a subject.

I. Introduction

Debugging is one of the most time-consuming and difficult aspects of software development [?], [?]. Recent years have seen a wide variety of research efforts devoted to easing the burden of debugging by *automatically localizing faults*. The most popular approaches, following the seminal work of Jones, Harrold, and Stasko [?], [?] use statistics of spectra [?] of failing and successful executions to score program entities according to how likely they are to be faulty. These spectrum-based approaches are popular in part because they have outperformed competing

approaches, and in part because they are highly efficient and easy to use — they typically only require the collection of coverage data and marking of tests as passing and failing, and thus are both computationally cheap and easy to fully automate. Many formulas have been proposed as potentially improving the accuracy of scores [?], [?], [?], [?], [?] over Tarantula, and a variety of complex and simple methods to modify the basic spectrum-based approach have been proposed [?], [?], [?].

Despite this large body of work and continuing interest, and the potentially high economic value of effective localization, there is no evidence of industrial adoption of spectrum-based localization [?] and there is recent concern about the long-term value of localization research. First, Parnin and Orso asked the core question: “Are automated debugging techniques actually helping programmers?” [?], and did not receive comforting answers. Second, it was shown that there is no single dominating “best” formula for localization [?]. Different formulas will perform best for different programs.

Parnin and Orso studied how actual programmers made use of localization techniques, including studying the effectiveness of localization using artificially high rank for faults [?]. Their conclusions included 1) a proposal to use absolute rank to measure effectiveness, because developers lose interest in localizations after a very few incorrect suggestions, and 2) focus on using richer information (e.g. the actual test cases) rather than just a raw localization in debugging aids. Combined with Yoo et al.’s establishment [?] that there is no truly optimal formula for localization, this suggests that the most valuable contributions to localization would be *formula-independent* methods that *potentially result in extremely large improvements in fault rank* rather than small, incremental average improvements in rank. This paper, therefore, argues that in many cases there *is* a simple, easily applied, improvement to localization that works with any formula (or other modification to the method we are aware of), has benefits to developers even if they ignore the localization, and often produces very large improvements in what we consider the most important measure of localization effectiveness, the absolute worst possible ranking of the faulty code.

A. Reduce Before You Localize

If failing test cases only executed faulty code, fault localization would not be an interesting research problem because fault localization would be the most trivial part of every debugging effort: localization would involve simply examining coverage of each failure. Of course, failing test cases usually execute a large number of non-faulty statements, which makes debugging very difficult. It is often impossible to produce a failing test case that executes only faulty code. However, it is also true that for any fault, there exist test cases that *execute as little failing code as possible*. Due to the way spectrum-based localizations work, reducing the amount of non-faulty code executed in failing test cases will almost always improve localization. Consider the Tarantula [?], [?] formula. Tarantula, like most spectrum-based approaches, determines how suspicious (likely to be faulty) a coverage entity e (typically a statement) is based on a few values computed over a test suite:

- $\text{passed}(e)$: # of tests covering e that pass
- $\text{failed}(e)$: # of tests covering e that fail
- totalpassed : the # of passing tests
- totalfailed : the # of failing tests

$$\text{suspiciousness}(e) = \frac{\frac{\text{failed}(e)}{\text{totalfailed}}}{\frac{\text{failed}(e)}{\text{totalfailed}} + \frac{\text{passed}(e)}{\text{totalpassed}}}$$

It is easy to see that if we lower $\text{failed}(e)$ for all non-faulty statements, while keeping everything else unchanged, the rank (in suspiciousness) of faulty statements will improve. If $\text{failed}(e)$ falls to 0 for all non-faulty statements, only faulty statements will be suspicious, a perfect localization. Reducing coverage of non-faulty statements in failing tests, then, is a potentially very effective and formula-independent approach to improving localizations.

Unfortunately, there is no method in the literature for reducing the amount of non-faulty code executed in test cases, to our knowledge.¹ However, there *is* a widely used method for reducing the size of failing test cases, delta-debugging.

Delta-debugging [?] (DD for short) is an algorithm for reducing the size of failing test cases. Delta-debugging algorithms have retained a common core since early proposals [?]: use a variation on binary search to remove individual components of a failing test case t to produce a *new* test case t_{1min} satisfying two properties: (1) t_{1min} fails and (2) removing any component from t_{1min} results in a test case that does not fail. Such a test case is called *1-minimal*. Because 1-minimal test cases are potentially

much larger than the smallest possible set of failing components, we say that delta-debugging *reduces* the size of a test case, rather than truly minimizing it. The precise details of delta-debugging and its variants can be complex; however, the family of delta debugging algorithms can generally be simply described. Ignoring caching and the details of the divide-and-conquer strategy for constructing candidate test cases, DD for a failing test case t_b proceeds by iterating the following two steps:

- 1) Construct the next candidate simplification of t_b , which we call t_c . Terminate if no t_c remain (t_b is 1-minimal).
- 2) Execute t_c by calling $\text{rtest}(t_c)$. If rtest returns \mathbf{X} (the test fails) then it is a simplification of t_b . Set $t_b = t_c$.

Delta-debugging reduces the size of a test case in terms of its components. Its purpose is to produce small test cases that are easier for humans to read and understand, and thus debug. In our long experience with delta-debugging [?], [?] and in recent work on variations and applications of delta-debugging [?], [?], [?], we noticed that in addition to reducing the “static” human-readable text of a test case, delta-debugging also almost always *reduces the code covered by a failing test case*, often by hundreds or thousands of lines [?]. The core proposal of this paper, therefore, is that *failing test cases should be, when possible, reduced with delta-debugging before they are used in spectrum-based fault localization: **reduce before you localize***. The original, unreduced, test cases should not be used, as they likely contain much irrelevant, non-faulty code that may mislead localization.

Given that delta-debugging is extremely valuable on its own merits in debugging (it is generally agreed that it is absolutely required for effective random testing, for example [?], [?], [?]) and fits well into Parnin and Orso’s proposal that debugging aids should focus on integrating all sources of information, including actual test cases, applying delta-debugging to failing test cases before localizing them is a low-cost, potentially high-value modification to spectrum-based fault localization. Even if reduction does not improve the localization, we show that under some reasonable assumptions it will not produce a *worse* localization, and at least the developer now has a set of smaller, easier-to-understand test cases to read. In fact, we believe *the only reason not to reduce test cases before localization, as a best practice, is when it is too onerous (or not possible) to set up delta-debugging for test cases*.

In this paper, we show that using delta-debugging to reduce failing test cases, when applicable, does usually produce improvements in fault ranking, using a variety of standard localization formula from the literature, and these improvements are often dramatic. In order to place our results on a firm empirical footing [?] we provide results over both SIR/Siemens [?] suite subjects studied in previous literature, a set of real faults from an industrial-strength random testing framework for the SpiderMonkey

¹Cause reduction [?], as we note later, could probably be used for this purpose, but might be expensive and has not been applied to this goal.

JavaScript engine [?], [?], and a variety of open source Java programs. Not only does reducing failing tests produce improvements; the improvements produced are often even better than those provided by optimally switching formula. While our results are not strong enough in terms of absolute ranking to ensure industrial adoption, they are a significant step towards an automated debugging ecosystem that can help real users debug real faults more quickly and more confidently.

II. Related Work

As discussed in the introduction, there is a very large body of work on spectrum-based fault localization (e.g. [?], [?], [?], [?], [?], [?], [?], [?], [?]), all of which informs our work. The most important motivational results for this paper are the investigation of Parnin and Orso [?] into the actual use of localizations for programmers, which inspired our evaluation methods, and the claim of Yoo et al. [?] that no single formula is best, which directed us to seek formula-independent improvements to localization. Our use of many programs and methods was inspired by the threats identified by Steimann et al. to empirical assessments of fault localizations [?]. While our assumptions are not theirs (e.g., that test suites have 100% statement coverage) the theoretical analysis of Xie et al. [?] supports some of our intuitions about which formulas are most useful for localization, and our belief that counting appearances of entities in failures (the counts reduction attempts to reduce) is most critical to localization.

The most similar actual proposed improvement to localization to ours is the entropy-based approach of Campos et al. [?] that uses EvoSuite [?] to improve test suites. The underlying approaches are quite different, but both aim to improve the spectra used in localization rather than change their interpretation. The primary advantage of their approach over ours is that it can be of use when test cases cannot be reduced; on the other hand, EvoSuite is probably considerably harder to apply for most developers than off-the-shelf delta-debugging, and the delta-debugged test cases themselves are highly useful debugging products. We also provide results for more localizations and many more programs (and types of programs). Examining the faults used in their work, we suspect some of the open source failing test cases would reduce to a trivial-to-debug test case covering almost no non-faulty code. Another similar approach (sharing the same novel aspect of changing the test cases examined rather than the scoring function) is that of Xuan and Monperrus, who propose a *purification* for test cases [?] that executes omitted assertions and uses dynamic slicing [?] to remove some code from failing test cases (parameterized by each assertion). In practice, their dynamic slicing is performing some of the work we relegate to delta-debugging, and some of the effectiveness of their method is presumably due to lowering coverage of non-faulty code in the same manner, though this is not

their explicit goal. Delta-debugging can remove code from unit tests that would be in any dynamic slice, since it does not have to respect any property but that the test case still fails. A core practical difference is that their approach *only* applies to unit tests of method calls (since the slicing is at the test level, not of the program tested), and that we believe delta-debugging tools are more widely used and easily applicable than slicing tools (e.g they are language-independent).

This paper also follows previous work on delta-debugging [?], [?], [?] and its value in debugging tasks. The most relevant recent work is the set of papers proposing that in addition to producing small test cases for humans to read, delta-debugging is a valuable tool in fully automated software engineering algorithms even if humans do not read the reduced tests: e.g., it is helpful for producing very fast regression suites [?], for improving coverage with symbolic execution [?], and for clustering/ranking test cases by the underlying fault involved [?].

III. Assumptions and Guarantees

In addition to being orthogonal to the spectrum-based localization formula used, test case reduction has a second major advantage independent of empirical results. Namely, under a set of assumptions that hold in many cases, *reduction can only improve, or leave unchanged, the effectiveness of localization*. All formulas for localization have some instances in which they diminish the effectiveness of localization compared to an alternative formula [?]. For example, while Tarantula is not the most effective formula in most studies, it is often best for an individual fault and test suite, and may be more robust to coincidentally correct tests than some other formulas [?]. Reducing failing tests before applying a formula, however, at worst leaves the effectiveness of localization unchanged, for most of the formulas in widespread use that we are aware of, under three assumptions:

- 1) all failing test cases used in the localization involve the same fault,
- 2) each failing test case reduces to a test case that fails due to the same fault as the original test case, and
- 3) reducing the input size (in components) also covers less code when the test executes.

The first assumption is probably the least likely to hold in some settings; however, it is also the assumption that is least relied upon. Reduction will only be harmful if some faulty code is removed from some failing test cases that executed it but where it was irrelevant to the fault. While this is possible, it is only harmful when the removal of non-faulty lines does not overcome this in producing an overall ranking of statements. That is, the localization will suffer when, for some reason, faulty lines that are not required

for a given failure are removed at a *higher rate* than non-faulty lines that are not required for that failure. The second assumption is a usual assumption of delta-debugging. Most delta-debugging setups are engineered with this goal in mind, often using some aspect of failure output or test case structure to keep “the same bug.” Observed “slippage” rates for faults seem to be fairly low [?], even with little mitigation. Note that in the setting where a program has a single fault, assumptions 1 and 2 always hold. In a single fault setting, the coverage vector for test cases for the faulty code (which must execute for failure to take place) will be the same before and after reduction of test cases.

Finally, it is certainly *possible* for delta-debugging based on input to increase coverage during execution. In our experience, this is moderately uncommon, but it is certainly possible with standard delta-debugging setups. In practice, this problem is very easily mitigated by adding as a second criteria for reduction that the test case not only still fail (in the same way) but that it execute no statements not executed by the original test case, an instance of cause reduction (generalized delta-debugging) [?]. In this paper we have not applied this additional requirement, to show that even using already-existing delta-debugging infrastructure, developers can greatly improve the effectiveness of fault localization by first reducing failures.

Given these assumptions, we now show that reduction is, at worst, harmless for most formulas. Recall that spectrum-based localizations rely on only a few values relevant to each entity e to be ranked in a localization: $\text{passed}(e)$, $\text{failed}(e)$, totalpassed , and totalfailed . Given assumptions 1-3 above, *for faulty statements all of these formula elements will be unchanged after delta-debugging*. For non-faulty statements, the only possible change is that $\text{failed}(e)$ may be lower than before failing test cases were reduced. Holding the other values constant, it is trivial to show that most formulas under consideration are (as we would expect), monotonically increasing in $\text{failed}(e)$. Therefore, after reduction, the suspiciousness scores for faulty statements are unchanged and the suspiciousness scores for non-faulty statements are either unchanged or lower. The rank of all faulty statements is therefore either unchanged or improved. There are many spectrum-based fault localization formulas in use. In our evaluation, we have used 3 well known examples, in addition to the basic Tarantula [?] formula, as representative:

Ochiai [?]:

$$\text{suspiciousness}(e) = \frac{\text{failed}(e)}{\sqrt{(\text{totalfailed})(\text{failed}(e) + \text{passed}(e))}}$$

Jaccard [?]:

$$\text{suspiciousness}(e) = \frac{\text{failed}(e)}{\text{failed}(e) + \text{totalfailed}}$$

SBI: [?], [?]

$$\text{suspiciousness}(e) = \frac{\text{failed}(e)}{\text{failed}(e) + \text{passed}(e)}$$

All these formulas are monotonically increasing in $\text{failed}(e)$.² Reduction can only theoretically improve fault localization or in worst case leave it unchanged if formula under consideration is monotonically increasing $\text{failed}(e)$. We chose these 4 formulas as they were used before to study effect of *test suite reduction on fault localization*, to study *effect of test case purification on fault localization* [?] [?]. One of the formula that we found not to be monotonically increasing in $\text{failed}(e)$ is AMPLE [?]. The guarantee that reduction can only improve fault localization effectiveness does not hold for such a case. We note that according to the theoretical analysis of Xie et al. [?], formulas making number of failed test cases containing e the primary determinant of rank are best under certain assumptions, which fits our expectations for the impact of reduction well.

A. Limitations of Reduction for Localization

Reduce before you localize only applies when you can actually reduce. First, some types of test inputs cannot be broken down into components and thus cannot be reduced: if a program takes as input a single integer, or fixed size vector of integers, we see no reasonable way to apply delta-debugging. Similarly, in some cases the difficulty of maintaining test case validity during reduction is too high — in such cases we also suspect that automated testing in general is hard to apply. Finally, in some cases delta-debugging may technically apply but be of little value because test cases are already nearly minimal, and little coverage is lost when reduction is applied. This happens more frequently than might be expected, for reasons discussed in Section IV-D, but we also show that this may be less of a problem than it first appears. In some cases, e.g. where the original test cases are very lengthy, reduction may be computationally expensive. However, when reduction is possible and significant, we suspect that its value for debugging is so high that it will be worth the effort, even if it is not helpful for localization.

IV. Experimental Results

A. Evaluation Measure

Because we aim to take into account the findings of Parnin and Orso [?], our evaluation of fault localizations is based on a *pessimistic absolute rank* of the highest ranked faulty statement. That is, for each set of suspiciousness metrics computed, our measure of effectiveness is the *worst possible position* at which the first faulty statement can be reached, when examining

²Proofs available on request, and to be posted on our web page as an appendix.

the code in suspiciousness-ranked order.³ For example, if ten statements all receive a suspiciousness score of 1.0 (the highest possible suspiciousness), and one of these is the fault, we assign this localization a rank of 10; an unlucky programmer might examine this statement last of the ten highest-ranked statements. Pessimistic rank nicely distinguishes this result from another localization that also places the bug at score 1.0, but gives twenty statements a 1.0 score. In our view, following Parnin and Orso [?], the most important goal of a localization is to direct the developer to a faulty statement as rapidly as possible, ignoring the size of the entire program or even of the faulty execution. We considered counting all statements in the same basic block as “the fault,” given that spectrum-based approaches using statement coverage cannot localize below this granularity, but rejected this due to the potential need to eventually compare results with more sophisticated methods that use state or other information to provide finer-grained localizations. In practice, due to basic block sizes and the fact that programmers won’t always be “unlucky” and read the faulty code last, pessimistic rank may not show all cases where a localization would be useful in practice, but its primary purpose here is comparison. We note that we did not provide statistical test information for our results because none of our experiments could be considered as sampled from some random distribution; given a suite and fault, results are not stochastic. Despite our efforts, the set of programs and faults used is not really a basis for precise statistical inferences. We note that performing such an analysis would, of course, show the improvements of reduction to be statistically significant.

B. SIR Programs

Our initial experiments use the Siemens/SIR [?], [?] suite programs studied in many previous papers on fault localization, in particular the classic evaluation of the Tarantula technique [?]. These subjects provide a large number of faults, reasonable-sized test suites, and have historically been used to evaluate localization methods.

Of the seven Siemens programs considered in the empirical evaluation of Tarantula, only one was unsuitable for delta debugging: TCAS takes as input a fixed-size vector of integers, and therefore its inputs cannot be easily decomposed. For the remainder of the programs, the input is easily considered as either 1) a sequence of characters or 2) a sequence of lines, when character-level delta debugging is not efficient (and so unlikely to be chosen by users in practice), which was required for the `tot_info` subject. In all cases, reduction took on average less than three seconds per failing test case, an essentially negligible computational cost.

We evaluated our proposal by 1) first computing the fault ranking for each version of each subject by the five

formulas then 2) performing the same computation, but using only reduced (by delta-debugging) versions of the failing tests. Reduction was performed using Zeller’s delta-debugging scripts, available on the web, and comparing the output of the original (correct) version of the program and the faulty version as a pass/fail oracle.

Figures 1 and 2 show the results⁴. The lighter shaded bars show the ranking of the fault, without any reduction. The darker bars show the ranking after delta-debugging all failures. The graphs are shown in log-scale due to the range of rankings involved. In many cases, reducing test cases before localizing improved the ranking of the fault by a factor of two or more. Results for individual subject vary: for `print_tokens`, the average ranking for faults, over all bugs and all formulas, is 59.5 without reduction, and 36.4 with reduction. The result is improved by reduction in 19 cases, remains the same in 15 cases, and is worse in 1 case. These numbers improve to 29.8 average ranking, 18 improvements, and 10 unchanged results if we ignore AMPLE scores. Table I shows similar data for all the SIR subjects.

For all but one subject, average scores were better after reduction. That one subject, `schedule` has worse results due to faulty version 7. This version, unlike all other SIR bugs, is not actually a single fault. Instead, two independent (but identical in text) changes to the code exist, and reduction takes many test cases that execute both faulty lines, and removes one of them. This results in both faulty lines getting worse ranks. In other words, our first assumption was violated. However, this is a somewhat unusual “multiple fault” case in that the two faults are identical incorrect statements, and in some sense interchangeable — in a more typical multi-fault setting, we expect that faults will tend to be independent and a test case for one fault will often not cover the other fault in the first place. That is, our belief is that the typical multi-fault setting has multiple faults, but test cases that involve only one fault. In this example, however, one fault executed in 19 of the 27 failing test cases, the other executed in 24, and *both* executed in 16: after reduction, of course, since either fault will cause a problem, both were covered less frequently, and no test executed both. The other non-AMPLE cases where the ranking was worse seem to all be due to decreasing input size leading to some increase in code coverage, due to the structure of the software. One simple solution for this problem would be to modify the delta-debugging scripts to reject reductions that cover code not covered by the original failing test case. Even without such a modification, if we ignore AMPLE results, improvement in fault rank was 1.3 times as common as no change in rank, and *nearly 10 times as common as worse rank for the fault*. In addition to the frequency of improvement of fault rank, it is also important to examine

³As usual since at least the work of Reneiris and Reiss[?], we consider reaching *any* faulty statement to be sufficient.

⁴Version 32 of `replace` is missing because the fault relies on the library definition of `isalnum`, and is no longer present on modern Linux systems, as we determined after communicating with Gregg Rothermel and the SIR maintainers.




Fig. 1. First Set of SIR Results

Subject	With AMPLE					Without AMPLE				
	Avg.	Avg. (DD)	#Better	#Same	#Worse	Avg.	Avg. (DD)	#Better	#Same	#Worse
print_tokens	59.7	36.4	19	15	1	59.7	29.7	18	10	0
print_tokens2	26.8	9.2	23	20	7	27.0	5.8	19	17	4
replace	26.5	24.6	38	98	19	24.5	21.8	37	76	11
schedule	13.1	22.7	18	16	11	7.7	14.3	18	14	4
schedule2	100.3	88.4	29	17	4	92.4	76.6	28	12	0
tot_info	34.1	26.05	82	24	9	29.7	17.7	75	16	1
Total			209	190	51			195	145	20

TABLE I. SIR Fault Rank Change Result Frequencies

the degree of improvement (or the opposite) provided by reduction. Table II shows, ignoring AMPLE, the min, max, and average for changes in rank. With the exception of the unusual multi-fault program, *schedule* v7, (which accounts for all worse ranks for *schedule*), the effect size when reduction improved rank was usually much larger than the effect size when it gave worse results. For *replace*, the subject with the most instances where reduction made fault ranking worse, we see that the effect size when reduction was harmful was much smaller than when reduction was helpful. Furthermore, when reduction helped, it often improved the ranking of the fault by more than *optimally* switching formula (the exceptions usually rising from a poor performance by AMPLE). That is, we can ask: if we compare taking the *worst* formula and applying reduction to improve the localization, how often

is this better than switching to the *best* localization formula for that subject and fault? Obviously applying reduction is more practical, since we don't know in advance which formula will perform best, until we know the correct result. By this comparison, if we ignore AMPLE, it was better to apply reduction than switch from *worst* to *best* formula in 36 cases over all SIR subjects. It was better to switch formula in only 22 cases (in the remaining 32 cases these two methods were tied).

C. SpiderMonkey JavaScript Engine

SpiderMonkey is the JavaScript Engine for Mozilla, an extremely widely used, security-critical interpreter/JIT compiler. SpiderMonkey has been the target of aggressive




Fig. 2. Second Set of SIR Results

Subject	Better			Worse		
	Min	Max	Avg.	Min	Max	Avg
print_tokens	6	85	44.6	N/A	N/A	N/A
print_tokens2	3	67	45.8	6	6	6.0
replace	1	30	9.6	1	3	1.4
schedule	1	15	4.8	85	75	81.3
schedule2	4	35	22.6	N/A	N/A	N/A
tot_info	1	81	14.7	3	3	3.0

TABLE II. SIR Fault Rank Change Effect Sizes

Bug#	Revision Fixed	#Failures	diff size
R60	1.16.2.1	1	115
R95	1.3.2.3.8	7	111
R115	1.4.8.1	4	592
R360	3.117.2.6	3	223
R880	3.17.2.14	28	272
R1172	3.208.2.63	150	214
R1294	3.241.2.1	405	80
R1543	3.36.16.1	146	169
R1561	3.37.2.1.4.1.2.2	2	31
R1873	3.50.2.29	1,041	56

TABLE III. Spidermonkey Bugs

random testing for many years now. A single fuzzing tool, jsfunfuzz [?], is responsible for identifying more than 1,700 previously unknown bugs in SpiderMonkey [?]. SpiderMonkey is (and was) very actively developed, with over 6,000 code commits in the period from 1/06 to 9/11 (nearly 4 commits/day). SpiderMonkey is thus ideal for evaluating how reduction aids localization when using a sophisticated random testing system, using the last public release of the jsfunfuzz tool [?], modified for swarm testing [?]. Using a set of faults in SpiderMonkey version 1.6 found with random testing in previous research [?], we show that reduction is essential for localization of bugs found using random testing, and that the use of reduction is even somewhat more important than the choice of localization formula. Note that this is an extremely challenging setting for spectrum-based localization: ran-

domly generated successful tests tend to cover a wide variety of behavior, some of it in a very shallow fashion, which makes distinguishing the signal of lines covered more frequently in failing tests from general noise (and particularly the noise of lines that are simply hard to cover and happen to appear in the failures) very difficult.

Figure 3 shows the change in rankings of the faulty code for 10 SpiderMonkey bugs (Table III). These bugs were taken from our PLDI 2013 data set [?]. Out of the 28 bugs studied in that paper we chose 10 random bugs for which, by hand, we could confirm the true set of faulty lines in the code commit. Each bug is identified by the revision number of the commit in which it was fixed: e.g., R0 maps to revision 1.10.4.1, the first commit




Fig. 3. SpiderMonkey Results

of Spidermonkey changes under consideration. Table III shows all bugs studied, the commit version fixing the bug, the number of failing test cases for that bug (# Failures), and the size (in lines) of the fixing commit’s `diff`. The faults under consideration here are clearly non-trivial (in fact, most fixes involved changes to multiple source files). For our localization we used the original and reduced test cases from PLDI 13 [?] plus a set of 720 randomly generated passing tests generated using the same method as in the original data set.

Across these 10 bugs, the average ranking for the first faulty line encountered was 1,550.7 without reduction, improving to 994.5 with reduction. Reduction improved the localization in 33 cases, with a minimum improvement of 1 ranking and a maximum improvement of 2,137 positions. The average improvement was 674 positions. The results were unchanged in 7 cases. It is important to note that even with such a challenging setting and real bugs, fault localization with reduction performs better than fault localization without reduction irrespective of formula being used, in worst case it performed as good. It was better to use reduction than to optimally (from worst to best) switch formula for 5 of the 10 bugs; it was better to switch formula in 4 cases, and in one case both methods gave the same result. While the results show that reduction was extremely effective in improving localization, it is also true that the localization was *still not very helpful* in many of these cases, considering absolute pessimistic rank as success criteria. Of course, SpiderMonkey 1.6 has over 80KLOC, and even reduced failing tests typically executed over 8,000 lines of code, so a “poor” localization may be useful in such a large fault search space. For 6 of the 10 bugs, all scores after reduction gave the fault a ranking of 128 or better. We also suspect that experienced developers could dismiss at least some of the lines in the rankings immediately, as we discuss next.

In order to expand on how reduction improves localiza-

tion, we examine in more detail the bug we call “R1543,” fixed by commit 3.36.16.1, and its Tarantula localization. In our experiments, there were 146 test cases that failed due to bug R1543. The relevant commit contains three kind of changes: (1) a rename refactoring, (2) a refactoring for readability, and (3) the actual fault fix. The fault is in `jsexn.c`, in the function `Exception`. Comparing diffs of R1542 and R1543 we were able to identify lines 583, 587, 589, 590 and 594 as the actual faulty code. Before reduction, these lines (all with suspiciousness 0.716605) had rankings from 191 on (ranking is arbitrary among buggy lines of same suspiciousness). After reduction, the rankings improved to start at 98. The suspiciousness values of faulty lines, as discussed in Section III, remained unchanged. However, many higher ranked non-faulty lines became less suspicious after being removed from failing tests. The average coverage for an unreduced failing test was 11,461 lines, which decreased to 8,349 (an average decrease of 3,112 lines) after reduction.

When we looked at the lines ranked before the faulty lines, we found that 70 lines in unreduced ranking and 56 lines in post-reduction rankings were assigned suspiciousness values of 1.0. These lines execute only in failing test cases. Our further investigation suggests that many of these lines are actually failure handling and reporting code. For example, 23 of the highly ranked lines for R1543 are part of the functions `my_ErrorReporter` and `Quit` in `js.c`, clearly failure handling and reporting code that would not slow an experienced SpiderMonkey developer in reaching the actual fault. We found such failure handling and reporting code in highly ranked lines for other bugs as well (finding them is fairly easy, as they have suspiciousness 1.0 by Tarantula). While many lines that are not such obviously non-faulty code remain for many SpiderMonkey bugs, this suggests a follow-on to Parnin and Orso’s experiments [?]: while they included some more experienced developers in their study, they did not investigate how developers highly skilled with *a*

Subject	Program Source			Test Suite
	#Classes.	#Methods	SLOC	#Test cases
Apache Commons Validator	64	578	6,033	434
JExel 1.0.0 beta 13	43	133	1,522	344
JAxe	167	1,078	12,462	2,138
JParser	115	178	3,046	647
Apache Commons CLI	23	208	2,667	364

TABLE IV. Open Source Subject Programs

particular code base use fault localization. We speculate that while the raw rankings for many SpiderMonkey bugs are not very high, some of these results could be more useful than is apparent, in the hands of expert developers. Such users are the most likely eventual industrial adopters of localization, we believe, for complex projects where debugging can be extremely difficult and time consuming and aggressive automated testing is used.⁵

D. Open Source Projects

The SIR results show that reduction is useful for localization in an idealized setting where suites are relatively complete and faults are chosen for certain properties, and the Spidermonkey results show its value in a particularly challenging setting (where passing tests have very wide-ranging coverage). How does reduction affect localization in a typical open source project?

We attempted to answer this question by picking five open source Java programs (shown in Table IV) and a set of identified and fixed bugs for those projects. We used patches that fixed the bugs to reproduce the bugs and then applied fault localization. Unfortunately, we found out that the dataset was not very helpful to effectively carry out fault localization. In almost every case, only one test case in the suite for the project caught the fault, and that test case was, essentially, already minimal. We examined a large number of bugs in the bug repository for Validator and CLI and confirmed our suspicion that this was very common: for most bugs, there was one, near-minimal, test case. The most test cases we saw for any bug was four, and those were, again, near-minimal. Combining the project suites and fixed bugs for these projects yielded a very uninteresting localization problem, from our point of view. Note that reduction was not harmful or expensive, it simply *made no real difference* for this data set. Why?

1) The Bugs Researchers Don't See: The following describes what we believe to be a common approach to building up the test suite and bug database for a project:

- 1) Initially, a small set of tests are written to check for basic functionality of the system.
- 2) Before checking in new code, developers run the established suite to make sure the code is largely correct. The developer may add tests for new functionality.
- 3) When users report a bug via the bug database that is not detected by the current suite, developers add a new test (or a small set of tests) to catch *that fault*. The fix for the bug is often checked into the code repository at the same time.

This is not an unreasonable approach, and roughly fits many development models, including test-driven development. The primary deviations are that some projects don't maintain a good bug database or add tests to catch all fixed bugs. This workflow has two unfortunate side effects for researchers, however. First, note that for bugs with entries in the bug database or fixes in the repository (the common source of faults to use in evaluating localization or testing methods), the test case(s) for the fault are likely to be already near-minimal, because *they are written specifically to catch that bug*. Second, if an existing, non-near-minimal test "accidentally" catches a bug during development, that bug *likely never makes it into the source history or bug database, because a developer will fix it before checking in the code*. The primary exception is when a test added to catch a newly discovered fault reveals older code that is also incorrect, for an unrelated reason. We speculate that this is a rare occurrence. Together, these two related issues mean that researchers using source history and the bug database to investigate fault detection or localization tend to see only a subset of the faults that have ever been present in some version of the software [?]. There is a "dark matter" of bugs encountered during development that were detected and fixed before they made it into the source repository. Moreover, we speculate that these bugs may be likely to have much more reducible test cases than those that are visible to researchers using traditional software history mining techniques, since no test case was written for each bug.

This is particularly problematic for evaluation of fault localization, because the bugs in the history may often not be the faults where a developer could have benefited most from spectrum-based localization. In many cases, we believe, the failing test case is produced based on user bug reports, and by the time a test case exists, the bug is often localized — the test case is in many cases derived from the fix, not the other way around! There are, of course, various localization techniques based on field reports, but these operate on less information than traditional spectrum-based localizations. We believe there *are* faults where spectrum-based localization would be helpful: the ones in the "dark matter," detected by a conscientious developer running the regression suite before a check-in. Some of these bugs are surely trivial to understand, but many others may require a fairly lengthy debugging session. The value of localization in this case, however, may be somewhat different than the

⁵We the authors have submitted Chrome JavaScript engine bugs found during testing research and seen that it may take months to understand and fix some complex failures.




Fig. 4. First Set of Open Source Results




Fig. 5. Second Set of Open Source Results

typical case. We expect that developers will often have a good idea what the faulty code is in these cases: it is probably code that just changed. However, in other cases the newly added code may only expose a fault in unchanged code. The consequences of faulty new code may also be hard to understand. In both of these cases, good localization could theoretically speed this “invisible” debugging process that (we speculate with some support in the literature [?]) often leaves no trace in the bug database, source repository, or even regression suite (the only trace might be lengthened time before the developer commits). The size and quality of “dark matter” bugs needs serious investigation, if fault localization research is to be responsive to real developer needs.

2) *Mutant-Based Evaluation*: There is no easy way to discover the missing bugs. However, we want to evaluate how reduction can improve fault localization. We therefore generated mutants for each of the open source projects as it is being done in very recent fault localization literature to simulate bugs.[cite] We do not claim that this necessarily

simulates small faults that might occur during development but never make it into the test suite or bug database, only that it is *plausible* the small changes in mutants resemble invisible bugs. A mutant [?] is a copy of the program with a single change to its source code; mutants are often used to substitute for real faults [?], including in recent fault localization work [?]. Our strategy was to create mutants in the same way as Xuan and Monperrus [?] in their own localization work, using 6 mutant operators. From each set of mutants generated, we selected 5 or 6 mutants at random that met the following criteria: (1) the mutant was killed by at least one test case and (2) the mutant generated no errors in JUnit test cases. A JUnit failure is caused by an unsatisfied assertion, but an error is caused by another kind of test failure, which may include some test setup or oracle problems. Using assertion failures only assures that we were remaining within the intent of the original tests.

Figures 4 and 5 show the results of applying reduction to these simulated faults. Taking all the open source projects and mutants together, we note that reduction improved fault ranking in 51 cases, left it unchanged in 55 cases, and made it worse in 2 cases. The average improvement was 17.62 ranking positions; the average negative effect size was 2 ranking positions. The best improvement was 100 rank positions. The average fault ranking without reduction was 37.64 and with reduction it improved to 29.36. However, compared to switching from worst to best formula, reduction did not perform as well here as with SIR or SpiderMonkey bugs: it was better to switch from worst to best formula in 14 cases, better to use reduction in 5 cases, and the methods were tied in 8 cases. Still, a simple, easily applicable method that is sometimes better than *knowing in advance which formula to use* and almost never harmful (while switching formula is quite often harmful, as the benefits of switching from worst to best formula show) is a highly useful addition to fault localization practice.

E. Threats to Validity

The primary threats to validity here are to external validity [?], despite our use of a larger number of faults and subjects than is usual in the literature. Our subjects are all C or Java programs, for example, and for the open source projects we used seeded rather than real faults. The all-too-frequent threats [?] of relying on results for a single formula or of making comparisons that are not across exactly duplicated subjects and code-bases, however, are absent from our study by design. To avoid construct threats, we developed independent experimental code-bases for some of the subjects, executed both, and compared results to cross-check the shared code base used for all subjects. Fortunately, most tasks here are straightforward (test execution, coverage collection, delta-debugging, and calculation of scores).

A second point (not strictly a threat) is that this paper

focuses on single-fault localization. Except for one case (where the SIR example has two separate faults), all our evaluations are over failures due to one fault. This is intentional, but limits the applicability of our results to multi-fault settings. The reasons for this choice are twofold: first, effective single-fault localization seems to be the minimal functionality required, and is easier to evaluate. Second, we suspect that the most likely real-world uses of localization will be in single fault settings, either due to the existence of only one detected fault (as in the “bugs during development” discussed above) or in aggressive automated testing of mature software. Even when there are multiple faults, it may be better to use techniques for clustering test cases by likely fault [?], [?], [?], [?] and then perform single-fault localization than to try to localize multiple faults at once.

V. Conclusions and Future Work

Our primary conclusion is that, when possible, anyone attempting to use spectrum-based fault localization should use delta-debugging to *reduce before localizing*. Across Siemens subjects, real Mozilla SpiderMonkey bugs, and mutants of a set of open source projects, reducing test cases before localizing was seldom harmful and in the cases where it caused harm the effect size was much smaller than in the cases where reduction was helpful. In most cases, reduction was helpful, and it was sometimes extremely effective, improving fault ranking by a factor of 2 (or more) and a very large absolute rank, sometimes hundreds of lines. This makes sense: if failing test cases only contained faulty code, fault localization would be trivial. Delta-debugging, by (usually) reducing the coverage of non-faulty code, approaches this ideal situation as best we know how at present. While delta-debugging is not a panacea for localization, in that it does not apply to some kinds of inputs and is sometimes not helpful, it often produces a very large improvement in localization effectiveness, quite often *more so than can be gained by switching from worst to best formula*.

In future work, we plan to examine whether adding constraints to avoid the infrequent cases where input size reduction *increases* coverage improves the value of reduction. There are a number of variations of this approach: e.g., simply disallowing coverage increases, or forcing reductions to actually reduce coverage as well as input size. In theory, mutation of input as well as reduction could be applied, to see if nearby inputs cover less code but still fail (however, this requires an algorithm other than delta-debugging, and introduces a very large state space). A second line of future work is to find a way to investigate actual “dark matter” bugs missing from the source histories and bug databases of projects. Our use of mutants to simulate such faults is somewhat plausible but clearly not satisfactory. One approach might be to use developer interviews or (more likely) captured IDE

history [?] to look into this type of fault, but this is resource-intensive and unlikely to be possible or remotely complete for many projects (e.g. open source projects have many developers, unlikely to all use an instrumented IDE). Simply confirming (or refuting) the development patterns we speculate may give rise to “dark matter” bugs would be a useful contribution to understanding the real problems of localization. More specific aspects of this question also arise: for example, what fraction of test cases in typical open source projects are specially crafted in response to a bug, rather than produced in hopes of detecting future bugs?

Our larger take-away message is that the lessons of Parnin and Orso [?] should be taken to heart: rather than seek incremental improvements in localization effectiveness, we need large improvements in fault rank, and need to exploit all sources of information, not just coverage vectors. Even when reduction does not assist localization, we believe that the reduced test cases are highly valuable debugging aids. Furthermore, because no single formula is “best” for all faults [?], there is much to be gained by devising aids to fault localization that apply to any formula and any type of spectrum. If automated fault localization is to be adopted in real-world settings, we need more than a competing set of ranking algorithms: we need a complete ecosystem for localization, debugging, and test understanding.